

Embedded Systems Summary

Andreas Biri, D-ITET

10.07.15

1. Introduction

Embedded Systems (ES): information processing systems embedded into a larger project

Cyber-physical system (CPS): must operate dependably, safely, securely, efficiently and in real-time

Characteristics of Embedded Systems (1-19)

- **dependable:** reliable, maintainable, available, safe
- **efficient:** energy, code, run-time, weight, cost
- **specialized:** dedicated towards certain application
- **real-time:** must meet constraints of environment
- not programmable by end-user
- fixed run-time requirements (additional power useless)
- criteria: cost, power consumption, predictability
- energy & temperature constraints (often independent)
- energy harvesting important (e.g. zero power systems)

Hard real-time constrain: not meeting that constrain could result in a catastrophe; answer arriving too late is wrong

Hybrid system: analog and digital system components

Reactive system: in continual interaction with environment executes at pace determined by environment

MPSoCs: Multiprocessor systems-on-a-chip (e.g. phone)

2. Software Introduction

Real-Time Systems (2-15)

ES are expected to finish tasks reliably within time bounds

Hard constrain: missing a deadline results in catastrophe often in safety-critical applications (aeronautics, brakes)

Soft constrain: missing deadline is undesirable but not fatal

Worst-Case Execution Time (WCET): upper bound on execution time of all tasks statically known

- difficult to calculate because of parallelism (branch prediction, speculation, pipelines) & caches

Best-Case Execution Time (BCET): lower bound for it

Programming Paradigms (2-25)

Time triggered approaches (2-26)

- periodic
- cyclic executive
- generic time-triggered scheduler
- no interrupts except by timer
- deterministic behaviour at run-time
- interaction with environment through polling

Summary

- + deterministic schedule (computed before run-time)
- + shared resources pose no problem


- external communication only via polling
- inflexible (no adaptation to environment)
- long processes have to be split into subtasks

Extension

- allow arbitrary interrupts (not deterministic anymore!)
- allow preemptable background processes

Simple Periodic Time-Triggered Scheduler

Timer interrupts regularly with period P (same for all processes)




- unpredictable starting times for later processes
- mutually exclusive, no sync required for communication

$$\sum_i WCET(T_i) < P$$

Time-Triggered Cyclic Executive Scheduler

processes may have different periods



- period P portioned into frames of length f
- terrible for long processes (need to be split)

Conditions

- Process executes at most once within a frame

$$f \leq p(k) \quad \forall k$$

- Period P is least common multiple of all periods $p(k)$
- Periods start and complete within a single frame:

$$f \geq WCET(k) \quad \forall k$$

- at least one frame boundary between release & deadline

$$2f - \gcd(p(k), f) \leq D(k) \quad \forall k$$

Generic Time-Triggered Scheduler

- precompute schedule a priori offline (if purely TT)

Task-Descriptor List (TDL): contains cyclic schedule for all activities, considering required precedence and mutual exclusion -> no explicit coordination at run-time necessary

Event triggered approaches (2-36)


- non-preemptive
- preemptive (stack policy, cooperative, multitasking)

Summary

- + dynamic & adaptive
- + can react to environment by receiving interrupts
- + guarantees can be given during run-time or even off-line
- problems with respect to timing
- shared resources have to be coordinated

Non-Preemptive Event-Triggered Scheduling

events are collected in a queue and cannot be preempted (cannot give guarantees regarding deadlines)



ISR: Interrupt service routine

- event associated with corresponding process
- events emitted by
 - a) external interrupts
 - b) processes themselves
- simple communication between processes
- buffer overflow if too many events are generated
- long processes prevent others from running (-> split)

Extension

- preemptable background process if event queue is empty
- timed events enter queue only after time interval elapsed

Preemptive Event-Triggered Scheduling

possible to preempt process, solves problem of long tasks

Stack-based: stack-based context mechanism of function calls (process = C-style function with own memory space)

- LIFO: restricts flexibility, bad if waiting for external event
- no mutual exclusion; shared resources must be protected (e.g. disable interrupt, semaphores)

Processes and CPU (2-43)

Process: unique execution of a program ("instance")

- has its own state (e.g. register values, memory stack)
- several copies of a program can run simultaneously

Activation record: copy of process state (includes registers)

Context switch: current CPU context goes, next comes

- context of current process is stored (registers, program counter, stack pointer)
- execution continues where other process left off

Co-operative Multitasking (2-45)

process allows context switch at `cswitch()` call

- + predictable where context switch can occur
- + less errors with use of shared resources
- bad programming can stall the system (doesn't yield)
- real-time behaviour at risk (if switch not possible)

Preemptive Multitasking (2-60)

Scheduler (OS)

- i) controls when context switches
- ii) determines which process runs next

Scheduler is called / switch enforced by:

- use of timers / timer interrupts
- hardware or software interrupts
- direct call to OS routines to switch context

3. Real-Time Models

Hard: missing its deadline has catastrophic consequences

Soft: meeting its deadline is desirable, but not critical

Schedule: assignment of tasks to the processor

- $\sigma(t) = 0$: processor is idle at time t
- $\sigma(t) = i$: processor is executing task i at time t

Feasible: tasks can be completed according to constraints

Schedulable: there exists at least one algorithm which can produce a feasible schedule

Schedule & Timing (3-5)

J_i / τ_i task / periodic task i

a_i / r_i arrival / release time (ready for execution)


C_i computation time (required CPU time)

d_i / D_i absolute / relative deadline $d_i \geq r_i + C_i$

s_i / f_i start / finishing time

T_i period (for periodic tasks)

Φ_i phase (start of periodic task)



Derived figures

Lateness $L_i = f_i - d_i$
delay of a task completion

Tardiness / exceeding time $E_i = \max(0, L_i)$
time exceeded after deadline

Laxity / slack time $X_i = d_i - a_i - C_i$
maximal time a task can be delayed on its activation to complete within deadline

Precedence Constrains: describes the interdependencies between tasks (“Which one has to be executed first?”)

Classification of Scheduling Algorithms (3-11)

Preemptive algorithm: running task can be interrupted at any time to assign the processor to another active task

Non-preemptive algorithm: once started, the task is executed until completion (no interruptions)

Static algorithm: scheduling decisions are based on fixed parameters, assigned to tasks before activation (offline)

Dynamic algorithm: scheduling decisions based on dynamic parameters that may change during system execution (e.g. CPU bursts, I/O waits)

Schedule metrics (3-13)

Optimal algorithm: minimizes given cost function

Heuristic algorithm: tends to find optimal schedule

Average response time

$$\bar{t}_r = \frac{1}{n} \sum_{i=1}^n (f_i - r_i)$$

Total completion time

$$t_c = \max(f_i) - \min(r_i)$$

Weighted sum of response time

$$t_w = \frac{\sum w_i (f_i - r_i)}{\sum w_i}$$

Maximum lateness

$$L_{max} = \max_i (f_i - d_i)$$

Number of late tasks

$$N_{late} = \sum_{i=1}^n miss(f_i), \quad miss(f_i) = \begin{cases} 0 & f_i \leq d_i \\ 1 & else \end{cases}$$

4. Periodic/Aperiodic Tasks

Aperiodic Tasks (4-3)

Equal arrival times & non-preemptive

- EDD (Jackson) for independent tasks
- LDF (Lawler) for dependent tasks

Arbitrary arrival times & preemptive

- EDF (Horn) for independent tasks
- EDF* (Chetto) for dependent tasks

Earliest Deadline Due (EDD) (4-4)

equal arrival times & non-preemptive : $O(n \log(n))$

Algorithm: Task with earliest deadline is processed first

Jackson’s rule: processing in order of non-decreasing deadlines is optimal with respect to minimizing the *maximum lateness*

Earliest Deadline First (EDF) (4-7)

arbitrary arrival times & preemptive : $O(n^2)$

Algorithm: Task with earliest deadline is processed first; if new task arrives with earlier deadline, current task is interrupted (just like EDD, but with recalculation)

Horn’s rule: executing the task with the earliest absolute deadline among the ready tasks at any time is optimal with respect to minimizing the *maximum lateness*

- $\sigma(t)$ task executing in the slice $[t, t + 1)$
- $E(t)$ ready task which has the earliest deadline
- $t_E(t)$ time at which the next slice of $E(t)$ is executed

Guarantee:

Worst case finishing time: $f_i = t + \sum_{k=1}^i c_k(t)$

EDF guarantee condition: $f_i \leq d_i \quad \forall i = 1, \dots, n$

A new tasks is accepted if the schedule remains feasible

Earliest Deadline First* (EDF*) (4-12)

determines a feasible schedule for tasks with precedence constrains if there exists one

Algorithm: Modify release times & deadlines, then EDF

Modification of release times:

1. Start at the top (*roots to leaves*)
2. Search the predecessor which takes the longest:

$$r_j^* = \max(r_j, \max(r_i^* + C_i : J_i \rightarrow J_j))$$

Modification of deadlines:

1. Start at the bottom (*leaves to roots*)
2. Search the successor which starts the earliest:

$$d_i^* = \min(d_i, \min(d_j^* - C_j : J_i \rightarrow J_j))$$

Latest Deadline First (LDF) (U2.2)

Non-preemptive scheduling for precedence constrains

Algorithm:

1. A precedence graph is constructed
2. *Leaves to roots:* Select task with latest deadline among all available tasks to be scheduled last
3. At runtime: tasks are extracted from head of the queue: first task inserted into queue will be executed last (*FIFO*)

Shortest Job First (SJF)

Minimizes average waiting time

Periodic Tasks (4-17)

Deadline *equals period*:

- Rate-monotonic (RM) for static priority
- EDF for dynamic priority

Deadline *smaller than period*:

- Deadline-monotonic (DM) for static priority
- EDF* for dynamic priority

Terminology

$\tau_{i,j}$ denotes the j -th **instance** of task i

$r_{i,j} / s_{i,j} / f_{i,j}$ **release / start / finishing time**

Φ_i **phase** of task i (release time of its first instance)

D_i **relative deadline** of task i (same for all instances)

T_i **period** with which the task is regularly activated

C_i **worst case execution time** (same for all instances)

Rate Monotonic Scheduling (RM) (4-22)

RM is *optimal* among all fixed-priority assignments, i.e. *no other fixed-priority algorithm can schedule* a task set which cannot be scheduled with RM

- static priority assignment (offline, as not changed)
- preemptive (by a task with higher priority)
- deadlines equals to the period ($C_i \leq D_i = T_i$)

Algorithm: tasks with higher request rate / shorter period will have higher priorities and interrupt lower ones

Critical instant: task is release simultaneously with all higher priority tasks / release creates largest response time

Schedulability analysis

Sufficient but not necessary (U : processor utilization factor):

$$U = \sum_{i=1}^n \frac{C_i}{T_i} \leq n (2^{1/n} - 1)$$

Sufficient and necessary: same as for DM

Deadline Monotonic Scheduling (DM) (4-34)

Deadlines may be smaller than the period:

$$C_i \leq D_i \leq T_i$$

Algorithm: tasks with smaller relative deadlines will have higher priorities and interrupt tasks with lower priority

Schedulability analysis

Sufficient but not necessary:

$$\sum_{i=1}^n \frac{C_i}{D_i} \leq n (2^{1/n} - 1)$$

Sufficient and necessary:

- worst-case demand when all tasks are released simultaneously (critical instances)
- **worst case interference** I_i for task i :

$$I_i = \sum_{j=1}^{i-1} \left\lceil \frac{t}{T_j} \right\rceil C_j$$

where tasks with $j < i$ have higher priority

- **Longest response time** $R_i = C_i + I_i$ (at critical instance)
- For **schedulability test**: find smallest R_i which satisfies

$$R_i = C_i + \sum_{j=1}^{i-1} \left\lceil \frac{R_i}{T_j} \right\rceil C_j \rightarrow R_i \leq D_i \quad \forall i$$

Earliest Deadline First (EDF) Scheduling (4-41)

Active task with earliest deadline has highest priority

- dynamic priority assignment
- preemptive
- $D_i \leq T_i$

Schedulability test ONLY for $D_i = T_i$

Necessary & sufficient: schedulable with EDF if and only if

$$\sum_{i=1}^n \frac{C_i}{T_i} = U \leq 1$$

U: average processor utilisation

Problem of Mixed Tasks Sets (4-47)


Periodic tasks: time-driven, execute regular critical control activities with hard timing constrains

Aperiodic tasks: event-driven; hard, soft or no real-time

Sporadic tasks: aperiodic task characterized by a minimum interarrival time (enables offline guarantee on constrains)

Background scheduling (4-48)

RM & EDF scheduling of periodic tasks: processing of aperiodic tasks in the background / when no periodic one



RM Polling Server (PS) (4-50)

Idea: Introduce artificial periodic task which services aperiodic requests as soon as possible

Function of polling server (PS): instantiated at regular intervals T_S and serves any pending aperiodic requests. If none, the process is suspended (time not preserved!)

Disadvantage: if an aperiodic request arrives just after the server is suspended, it must wait for next polling period

Schedulability analysis: just like RM, suff. but not necessary

$$\frac{C_S}{T_S} + \sum_{i=1}^n \frac{C_i}{T_i} \leq (n+1) (2^{1/(n+1)} - 1)$$

Sufficient if aperiodic task finishes before a new arrives

$$\left(1 + \left\lceil \frac{C_a}{C_S} \right\rceil \right) T_S \leq D_a$$

EDF – Total Bandwidth Server (4-55)

When k -th aperiodic request arrives at time $t = r_k$, it receives a deadline

$$d_k = \max(r_k, d_{k-1}) + \frac{C_k}{U_s}$$

$U_s = 1 - U_p$: server utilization factor / bandwidth

Once a deadline is assigned, the request is inserted into the ready queue as any other periodic instance

Schedulability test: necessary & sufficient

$$U_p + U_s \leq 1$$

5. Resource Sharing


Common resources: data structures, variables, main memory area, file, set of registers, I/O unit

Critical section: piece of code, in which access to shared resources requires **mutual exclusion**

blocked: task waits for an exclusive resource to be freed

holds: task is in possession of said resource

free: exclusive resource after leaving critical section



Semaphores (5-5)

S_i protects each exclusive resource R_i

$wait(S_i)$: start of critical section, requests entrance

$signal(S_i)$: end of critical section, frees resource

Priority Inversion (5-7)

- low-priority task holds resource which prevents high-priority task from running
- meanwhile, a medium-priority task can preempt the low-priority task and execute with the high-priority blocked

“Solution”: disallow preemption in critical sections

- unnecessary blocking of unrelated tasks with higher prio

Priority Inheritance Protocol (PIP) (5-10)

assume priority of highest blocked task in critical section

P_i : nominal priority

$p_i \geq P_i$: active priority

Direct Blocking: lower-priority task blocks higher task

Push-through Blocking: medium-priority task is blocked by low-priority task which has inherited a higher priority

6. Real-Time OS

Deficits of Desktop OS

- monolithic kernel too feature rich, takes too much space
- not: modular, fault-tolerant, configurable, modifiable
- not power optimized
- timing uncertainty too large

Advantages of Embedded OS

- OS can be fitted to each individual need: remove unused functions, conditions compilation depending on hardware, replace dynamic data by static data, advanced compiling
- improved predictability (everything through scheduler)
- interrupts can be employed by all processes
- software tested and considered reliable (no protection)

Real-Time OS (6-10)

Requirements

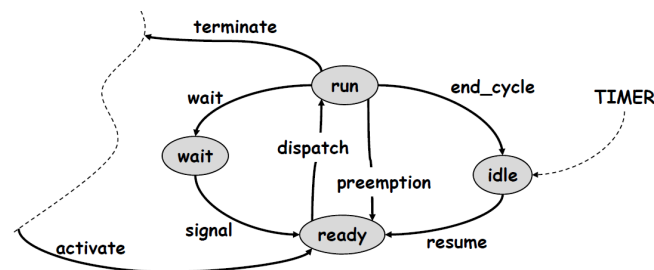
- **predictability of time-behaviour**
 - upper bound on the execution time of tasks
 - almost all activities controlled by scheduler
- **management of timing and scheduling**
 - inclusion of deadlines
 - OS must provide precise time services
- **speed**

Main functionality of RTOS-Kernels (6-13)

Process management (6-13)

- **execution of quasi-parallel tasks**
 - maintain process states & process queues
 - preemptive scheduling (fast context switch)
 - quick interrupt handling
- **CPU scheduling**: guarantee deadlines & fairness
- **Process synchronization** (semaphores, mutual exclusion)
- **Inter-process communication** (buffering)
- **real-time clock** for internal time reference

Process States (6-15)



run: starts executing on the processor

ready: ready to execute but not assigned yet

wait: task is waiting for a semaphore for access

idle: completed execution & waiting for next period

Threads (6-17)

A basic unit of CPU utilization, similar to a process

- typically shared: memory
- typically owned: registers, stack

Process: difficult to communicate, think they are alone

Thread: communicate via memory, knows there are others
multiple threads for each distinct activity of process

- faster to switch between threads (no major OS operation)
- **Thread Control Block (TCB)** stores information

Communication Mechanisms (6-20)

Problem: the use of shared memory for message passing may cause priority inversion and blocking

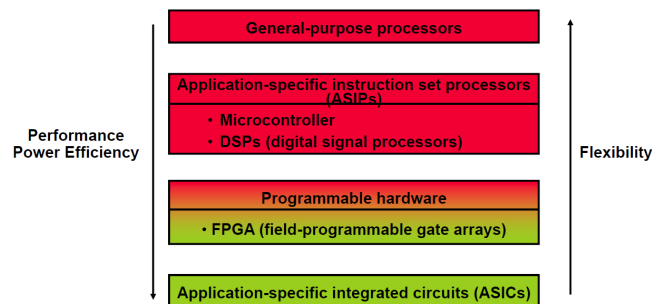
Synchronous communication ("rendez-vous")

- when communicating, they have to wait for each other
- causes problems for maximum blocking time
- in static RT environments solved offline by transforming synchronous interactions into precedence constraints

Asynchronous communication ("mailbox")

- sender deposits message into channel, receiver reads
- done by shared memory buffer, FIFO queue (fixed size)

7. System Components



General-purpose Processors (7-7)

- **high performance**
 - highly optimized circuits and technology
 - use of parallelism (pipelining, predictions)
 - complex memory hierarchy
- **not suited for real-time applications** as highly unpredictable execution times due to intensive resource sharing and dynamic decisions
- **good average performance** for large application mix
- **high power consumption**
- **Multicore Processors**
 - higher execution performance through parallelism
 - useful in high-performance embedded systems
 - interference on shared resources (buses, cache etc.)

System Specialization (7-13)

Specialization is main difference between embedded systems and general purpose high-volume microprocessors

- Specialization **should respect flexibility**
 - systems should cover a class of applications
 - required for later changes & debugging
- **System analysis required** for identification of application properties which benefit from specialization

Application-Specific Instruction Sets (7-22)

Microcontrollers / Control Dominated Systems

- Reactive systems with event driven behavior
- system description: *Finite State Machines* or *Petri Nets*

Microcontrollers **connect interfaces** (no computation)

- support process scheduling and synchronization
- preemption (interrupt), context switch
- short latency times
- low power consumption
- peripheral units often integrated (timer, buses, AD/DA-C)
- suited for real-time applications

Digital Signal Processors (DSPs) /

Data Dominated Systems (7-26)

- Streaming-oriented systems with periodic behaviour
- input description: flow graphs

DSPs are for **computation** (signal processing, controlling)

- optimized for data-flow, only simple control-flow
- parallel hardware units (VLIW), specialized instruction set
- high data throughput, zero-overhead loops
- suited for real-time applications

Very Long Instruction Word (VLIW): detection of possible parallelism by compiler, combine multiple functional units

Field Programmable Gate Array (FPGA) (7-34)

- "program hardware by software"
- **granularity of logic units:** gate, tables, memory, blocks
- **communication network:** crossbar, hierarchical mesh
- **reconfiguration:** dynamically adjustable at runtime

Application-Specific Circuits (ASICs) (7-41)

- **custom-designed circuits** for mass production
- long design times, lack of flexibility, high design costs

System-on-Chip (SoC) (7-43)

8. Communication

Requirements

- **performance** (bandwidth & latency, real-time)
- **efficiency** (cost, low power)
- **robustness** (fault tolerance, maintainability, safety)

Time Multiplex Communication (8-5)

Random Access (8-6)

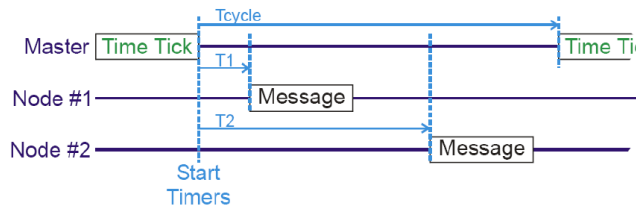
No access control, requires low medium utilization

Improved variant: *slotted* random access

TDMA (Time Division Multiple Access) (8-7)

Communication in statically allocated time slots

- synchronization among all nodes necessary
- master node sends out a synchronization frame



CSMA/CD (Carrier Sense MA / Collision Detection) (8-8)

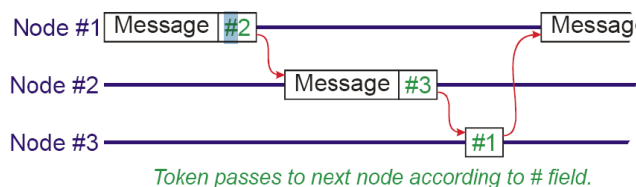
Try to avoid and detect collisions

- before transmitting, check whether channel is idle
- if collision detected, back off / wait
- repeated collisions result in increasing backoff times

Token Protocol (Token Ring) (8-9)

Token value determines which node is transmitting


- only the token holder may transmit



CSMA/ Collision Avoidance – Flexible TDMA (FTDMA) (8-11)

Reserve s slots for n nodes ; if slot is used, it becomes **slice**

- node start transmitting message only during assigned slot
- $s = n$: no collision ; $s \leq n$: statistical collision avoidance



CSMA/ Collision Resolution (CSMA/CR) (8-12)

Each node is assigned a unique identification number

- all nodes wishing to transmit send a binary signal based on their identification number; if node detects a dominant state while transmitting a passive one, it drops out
- node with the lowest identification value wins

Flex Ray (8-14)

Operation principle: Cycle is subdivided into static and dynamic segment. Static segment bases on fixed allocation of time slots, dynamic segment for ad-hoc communication

Static Segment: TDMA All static slots are the same length and are repeated every communication cycle

Dynamic Segment: Flexible TDMA minislot is opportunity to send a message; if not sent, minislot elapses unused

Bluetooth (Frequency Multiplex Communication) (8-20)

Design goals

- small size, low cost, low energy
- secure & robust transmission (interference with WLAN)

Technical Data

- 2.4 GHz (spectral bandwidth 79 MHz)
- 10-100m transmission range, 1 Mbit/s bandwidth
- simultaneous transmission of multimedia & data
- ad hoc network (de-centralized, dynamic connections)

Frequency Hopping

- transmitter jumps between frequencies: **1600 hops/s**
- 79 channels, ordering by pseudo-random sequence
- Frequency range: **$(2402 + k) \text{ MHz}$, $k = 0 \dots 78$**
- Data transmission in time window of **625 μs**
- Each packet transmitted on a different frequency

Network Topologies (8-24)

Ad-hoc networks

- all nodes are potentially mobile
- dynamic emergence of connections
- hierarchical structure (scatternet) of small nets (piconet)

Piconet

- contains **1 master and maximally 7 slaves**
- all nodes inside use the **same frequency hopping scheme** (determined by device address of master BD_ADDR)
- connections exist :
 - one-to-one
 - master and all slaves (broadcast)

Scatternet

- formed by several piconets with overlapping nodes
- node can be master in at most one and slave in other nets

Addressing (8-30)

Packet format

- Access Code / BD_ADDR : 82bits, identifies packets
- Header / AM_ADDR : 54bits, identifies connection
- Payload : 0 – 2745 bits

Bluetooth Device Address BD_ADDR : 48 Bits, unique

Active Member Address AM_ADDR :

- 3 bits for maximally 7 active slaves in piconet
- Address "Null" is broadcast to all slaves

Parked Member Address PM_ADDR : 8 bits

- in low power state: waiting for communication

Connection Types (8-31)

Synchronous Connection-Oriented (SCO)


- point-to-point full duplex between master & slaves
- master reserves slots for transmission regularly

Asynchronous Connection-Less (ACL)

- asynchronous service, no slot reservation
- master transmits spontaneous, slaves answer next

Frequency Hopping / Time Multiplexing (8-32)

- packet of the master is followed by a slave packet
- after each packet, channel / frequency is switched



- master can only start sending in even slot numbers
- packets have length of 1, 3 or 5 slots (same frequency)

Modes and States (8-35)

Modes of operation

Inquiry: master identifies addresses of neighbors

Page: master attempts connection with slave

Connected: connection is established

States in connection mode

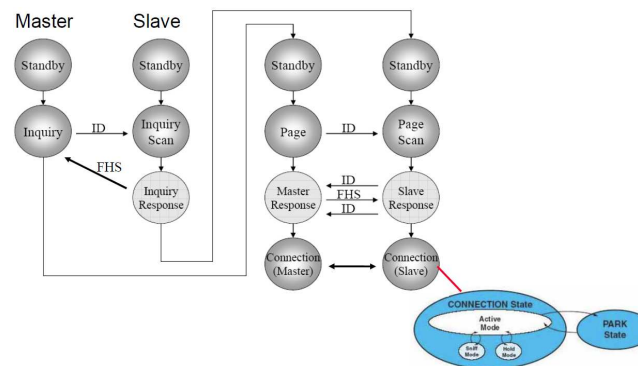
- *active* active in connection to master
- *hold* does not process data packets
- *sniff* awakens at regular intervals checks whether there are packets
- *park* passive, only synchronized

Synchronization in Connection Mode: channel sequence & phase of a piconet is determined (by BD_ADDR) of master

Synchronization in Page Mode: 3-way-handshake to synchronize between master and slave; prerequisite for establishing a connection

1. *Page:* master transmits own & slave address
2. *Page scan:* slave listens
3. *Master page response:* slave answers with own address
4. *Slave page response:* master sends FHS-packet, which includes channel sequence & phase of piconet

From Standby to Connection (8-40)



Protocol Hierarchy (8-44)

Baseband specification: defines packet formats, physical & logical channels, error correction, synchronization and modes of operations

Audio specification: defines coding & decoding

Link manager (LM): authentication & encryption, management, connection initiation, transitions

Host controller interface (HCI): interface host - node

Link layer control & adaption layer (L2CAP): interface for data communication

RFCOMM: simple transport protocol for serial connection

9. Low Power Design

Power is most important constrain in Embedded Systems

Power and Energy (9-9)

$$E = \int P(t) dt$$

Minimizing power consumption is important for

- design of the power supply & voltage regulators
- the dimensioning of interconnect
- cooling (decrease temporary heating)

Minimizing energy consumption is important due to

- restricted availability of energy (mobile systems)
- limited battery capacities & long lifetimes needed
- very high costs of energy (solar panels, in space)

Power Consumption of CMOS Processor (9-12)

Dynamic power consumption: charging & discharging C_L

Short circuit power consumption: switching causes shorts

Leakage: leaking diodes & translators, causes static current

Power $P \sim \alpha C_L V_{dd}^2 f$

Energy $E \sim \alpha C_L V_{dd}^2 f t = \alpha C_L V_{dd}^2 (\#cycles)$

Delay $\tau \sim C_L \frac{V_{dd}}{(V_{dd} - V_T)^2}$

V_{dd} : supply voltage

$V_T \ll V_{dd}$: threshold voltage

α : switching activity (= 1 : switch every cycle)

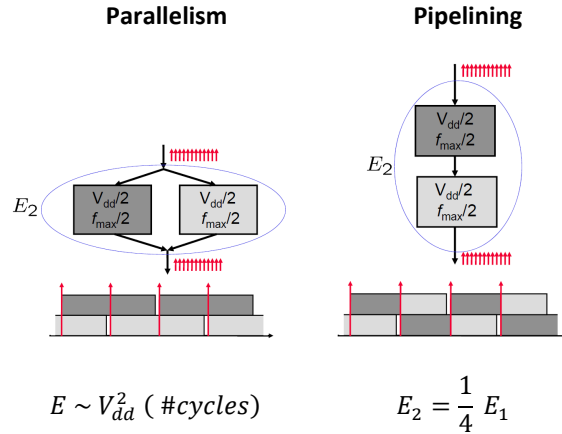
C_L : load capacity

$f \sim \frac{1}{\tau} \sim \frac{V_{dd}}{C_L}$: clock frequency

Basic Techniques (9-17)

Power Supply Gating: minimize static power consumption (leakage) by cutting off power supply while unit inactive

Parallelism (9-18)



VLIW Architectures (9-22)

Large degree of parallelism & many computational units:

- explicit parallelism (parallel instruction set) by compiler
- parallelization through hardware (difficult & expensive)

Translation of instruction set

- done with optimized compiler (no compatibility)
- on processor with decoder (translation in HW)
- on processor with dynamic compiler in SW (Transmeta)

Dynamic Voltage Scaling (DVS) (9-26)

Adapt voltage & frequency to situation to save energy

Optimal Strategy: running at a constant frequency/voltage minimizes energy consumption for dynamic voltage scaling

- if a task finishes on deadline, the chosen frequency (voltage) is optimal in terms of energy efficiency
- if only discrete voltage levels, choose directly above and below the ideal voltage to minimize energy consumption

YDS Algorithm for Offline Scheduling (9-36)

Schedule without missing deadlines & minimal energy
 $O(N^3)$, N : number of tasks in V

Intensity G in time interval $[z, z']$: average accumulated execution time of all tasks inside the interval

$$V'([z, z']) = \{v_i \in V : z \leq a_i < d_i \leq z'\}$$

$$G([z, z']) = \sum_{v_i \in V'} c_i / (z' - z)$$

1. Find **critical interval** (i.e. interval with highest intensity) and schedule tasks inside with EDF

$$C_{eff} = \frac{C_{tot}}{G}, \quad f = G * f_{nominal}$$

2. Adjust arrival times and deadlines by excluding interval
3. Run algorithm for revised input and put pieces together

Online algorithm: run algorithm with known tasks, if new ones arrive, update schedule; maximally uses 27 times the minimal energy consumption of optimal offline solution

Dynamic Power Management (DPM) (9-46)

DPM tries to assign optimal power saving states

RUN: operational

IDLE: SW routine may stop the CPU when not in use, while monitoring interrupts

SLEEP: shutdown of on-chip activity

DVS Critical frequency (voltage): running at any frequency (voltage) below is not worthwhile for execution

Procrastination Schedule: execute only voltages higher or equal to the critical voltage (round up lower ones)
 - procrastinate task execution & sleep as long as possible


10. Architecture Models

Dependence graph (DG) (10-4)

directed graph $G = (V, E)$, $E \subseteq V \times V$

$(v_1, v_2) \in E$:
 - v_1 (immediate) predecessor of v_2
 - v_2 (immediate) successor of v_1

- nodes represent tasks, edges represent relations
- describes order relations for execution of single tasks
- represents parallelism, not branches in control flow



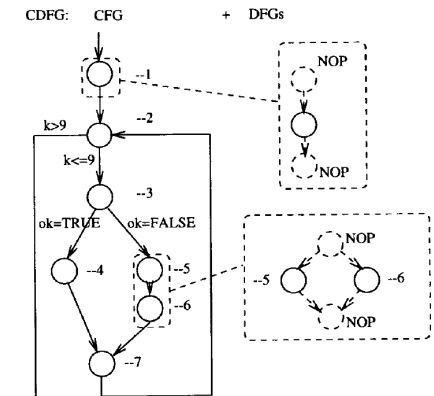
Control-Data Flow Graph (CDFG) (10-8)

Description of control structures & data dependencies
 - combines control flow & dependence representation

Control Flow Graph: finite state machine which represents the sequential control flow of the program (i.e. branches)
 - operations within state are written as dependence graph

Dependence Graph/ Data Flow Graph (DFG):

- NOP operations represent start and end point (polar)



Scheduling with resource constraints (11-34)

Minimal latency is defined as

$$L = \min\{\tau(v_n) : (\tau(v_j) - \tau(v_i) \geq w(v_i, \beta(v_i)) \forall (v_i, v_j) \in E_S) \wedge (|\{v_s : \beta(v_s) = v_t \wedge \tau(v_s) \leq t < \tau(v_s) + w(v_s, v_t)\}| \leq \alpha(v_t) \forall v_t \in V_T, \forall 1 \leq t \leq L_{max})\}$$

List Scheduling (11-36)

- static priority, which denotes urgency of being scheduled (e.g. higher priority, the further still away from end)
- algorithm schedules one time after the other and chooses from the tasks with top-priority
- heuristic algorithm, doesn't guarantee optimal scheduling

```
LIST( $G_S(V_S, E_S), G_R(V_R, E_R), \alpha, \beta, priorities$ ){
   $t = 1$ ;
  REPEAT {
    FORALL  $v_k \in V_T$  {
      determine candidates to be scheduled  $U_k$ ;
      determine running operations  $T_k$ ;
      choose  $S_k \subseteq U_k$  with maximal priority
      and  $|S_k| + |T_k| \leq \alpha(v_k)$ ;
       $\tau(v_i) = t \forall v_i \in S_k$ ; }
     $t = t + 1$ ;
  } UNTIL ( $v_n$  planned)
  RETURN ( $\tau$ ); }
```

Integer Linear Programming (11-42)

- yields optimal solution, as based on exact description
- binding already determined (know duration)
- know earliest & latest starting times from ASAP / ALAP

1. Minimize:

$$\tau(v_n) - \tau(v_0) = L$$

2. Decision variables x binary:

$$x_{i,t} \in \{0, 1\}, \forall v_i \in V_S, \forall t : l_i \leq t \leq h_i$$

3. Exactly one variable $x_{i,t}$ for all t has the value 1:

$$\sum_{t=l_i}^{h_i} x_{i,t} = 1, \forall v_i \in V_S$$

Each task can only have one starting point.

4. If $x_{i,t} = 1$ then the operation v_i starts at time t , i.e. $\tau(v_i) = t$.

$$\sum_{t=l_i}^{h_i} t * x_{i,t} = \tau(v_i), \forall v_i \in V_S$$

5. Precedence constraints are satisfied:

$$\tau(v_j) - \tau(v_i) \geq w(v_i), \forall (v_i, v_j) \in E_S$$

6. Resource constraints are not violated:

$$\sum_{\substack{v_i, v_k \in E_R \\ \forall i: (v_i, v_k) \in E_R}} \sum_{p'=\max\{0, t-h_i\}}^{\min\{w(v_i)-1, t-l_i\}} x_{i,t-p'} \leq \alpha(v_k) \\ \forall v_k \in V_T, \forall t : 1 \leq t \leq \max\{h_i : v_i \in V_S\}$$

Iterative Algorithms (11-49)

Consist of a set of indexed equations that are evaluated for all values of an index variable (e.g. signal flow graphs, marked graphs)

Representation of iterative algorithms

- one indexed equation with constant index dependencies
- equivalent set of indexed equations
- extended sequence graph denoting the displacements
- marked graph denoting displacement as data in queue
- signal flow graph (with displacement z^{-1})
- loop program

Definitions

Iteration: set of all operations necessary for computation

Iteration interval P: time distance between two iterations

Throughput 1/P: iterations per time unit

Latency L: maximal time distance between starting and finishing times of operations belonging to one iteration

Implementation Principles

- **Simple possibility:** edges with $d_{ij} > 0$ are removed and the resulting simple sequence graph solved traditionally

- **functional pipelining:** Simultaneous execution of data sets belonging to different iterations. Successive iterations overlap and a higher throughput is obtained

Solving the synthesis problem using Integer Linear Programming: (11-56)

- use extended sequence graph
- calculate upper and lower bounds as well as P
- replace equations (5) and (6) for ILPs

Dynamic Voltage Scaling (DVS) (11-60)

We can optimize the energy in case of DVS

- there are $|K|$ different voltage levels
- task $v_i \in V_S$ can use one of the execution times $w_k(v_i)$ and corresponding energy $e_k(v_i)$

1. Minimize:

$$\sum_{k \in K} \sum_{v_i \in V_S} y_{ik} * e_k(v_i)$$

Sums up all individual energies of operations.

2. Decision variables y_{ik} binary:

$$y_{ik} \in \{0, 1\}, \forall v_i \in V_S, k \in K$$

3. Exactly one implementation (voltage) $k \in K$ is chosen for each operation v_i :

$$\sum_{k \in K} y_{ik} = 1, \forall v_i \in V_S$$

4. Precedence constraints, where the actual execution time is selected from the set of all available ones:


$$\tau(v_j) - \tau(v_i) \geq \sum_{k \in K} y_{ik} * w_k(v_i), \forall (v_i, v_j) \in E_S$$

5. Guarantees deadlines:

$$\tau(v_i) + \sum_{k \in K} y_{ik} * w_k(v_i) \leq d(v_i), \forall v_i \in V_S$$

12. Various

Petri Nets (2-47)



- bipartite graph consisting of places and transitions
- data and control represented by moving tokens

Firing: enabled if at least one token in every input place
Remove one from each input and put one to each output

NutOS & Programming Practice (2-50)

Creating a thread

```
THREAD(my_thread, arg) {
    for (;;) {
        // do something
    }
}
```

a thread looks like a function that never returns

the thread is put into life


```
int main(void) {
    if (0 == NutThreadCreate("My Thread", my_thread, 0, 192)) {
        // Creating the thread failed
    }
    for (;;) {
        // do something
    }
}
```

stack size

sleep queue



thick lines: threads
thin lines: control



Terminating a thread

```
THREAD(my_thread, arg) {
    for (;;) {
        // do something
        if (some condition)
            NutThreadExit()
    }
}
```

can only kill itself

Yield access to another thread / set priority

```
THREAD(my_thread, arg) {
    for (;;) {
        NutThreadSetPriority(20);
        // do something
    }
}
```

Sleep

```
THREAD(my_thread, arg) {
    for (;;) {
        // do something
        NutSleep(1000);
    }
}
```

Posting & waiting for events (2-57)

```
#include <sys/event.h>

HANDLE my_event;

THREAD(thread_A, arg) {
    for (;;) {
        // some code
        NutEventWait(&my_event, NUT_WAIT_INFINITE);
        // some code
    }
}

THREAD(thread_B, arg) {
    for (;;) {
        // some code
        NutEventPost(&my_event);
        // some code
    }
}
```

wait for event, but only limited time

post event

Laboratory